

Importing Libraries and dataset

```


import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

```

```

df=pd.read_csv('/content/census-income .csv')
df

```



| | age | workclass | fnlwgt | education | education-num | marital-status | occupation | relationship | race | sex | capital-gain | capital-loss | hours-per-week | native-country |
|-------|-----|------------------|--------|------------|---------------|--------------------|-------------------|---------------|-------|--------|--------------|--------------|----------------|----------------|
| 0 | 39 | State-gov | 77516 | Bachelors | 13 | Never-married | Adm-clerical | Not-in-family | White | Male | 2174 | 0 | 40 | Unit Sta |
| 1 | 50 | Self-emp-not-inc | 83311 | Bachelors | 13 | Married-civ-spouse | Exec-managerial | Husband | White | Male | 0 | 0 | 13 | Unit Sta |
| 2 | 38 | Private | 215646 | HS-grad | 9 | Divorced | Handlers-cleaners | Not-in-family | White | Male | 0 | 0 | 40 | Unit Sta |
| 3 | 53 | Private | 234721 | 11th | 7 | Married-civ-spouse | Handlers-cleaners | Husband | Black | Male | 0 | 0 | 40 | Unit Sta |
| 4 | 28 | Private | 338409 | Bachelors | 13 | Married-civ-spouse | Prof-specialty | Wife | Black | Female | 0 | 0 | 40 | Ct |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 32556 | 27 | Private | 257302 | Assoc-acdm | 12 | Married-civ-spouse | Tech-support | Wife | White | Female | 0 | 0 | 38 | Unit Sta |
| 32557 | 40 | Private | 154374 | HS-grad | 9 | Married-civ-spouse | Machine-op-inspct | Husband | White | Male | 0 | 0 | 40 | Unit Sta |
| 32558 | 58 | Private | 151910 | HS-grad | 9 | Widowed | Adm-clerical | Unmarried | White | Female | 0 | 0 | 40 | Unit Sta |
| 32559 | 22 | Private | 201490 | HS-grad | 9 | Never-married | Adm-clerical | Own-child | White | Male | 0 | 0 | 20 | Unit Sta |
| 32560 | 52 | Self-emp-inc | 287927 | HS-grad | 9 | Married-civ-spouse | Exec-managerial | Wife | White | Female | 15024 | 0 | 40 | Unit Sta |

32561 rows × 15 columns

Next steps:


[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

EDA

```

df.isnull().sum().sum()

```




```
np.int64(0)
```

```

df.duplicated().sum()

```



```
np.int64(24)
```

```

df.drop_duplicates(inplace=True)
df.duplicated().sum()

```



```
np.int64(0)
```

```




df.info()




```

```
↔ <class 'pandas.core.frame.DataFrame'>
Index: 32537 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32537 non-null  int64
1   workclass              32537 non-null  object
2   fnlwgt                 32537 non-null  int64
3   education              32537 non-null  object
4   education-num          32537 non-null  int64
5   marital-status         32537 non-null  object
6   occupation             32537 non-null  object
7   relationship           32537 non-null  object
8   race                   32537 non-null  object
9   sex                    32537 non-null  object
10  capital-gain            32537 non-null  int64
11  capital-loss            32537 non-null  int64
12  hours-per-week          32537 non-null  int64
13  native-country         32537 non-null  object
14  annual_income          32537 non-null  object
dtypes: int64(6), object(9)
memory usage: 4.0+ MB
```

✓ Outliers





```
for x in df.columns:
    if df[x].dtypes!='object':
        sns.boxplot(df[x])
        plt.xlabel(x)
        plt.show()
```



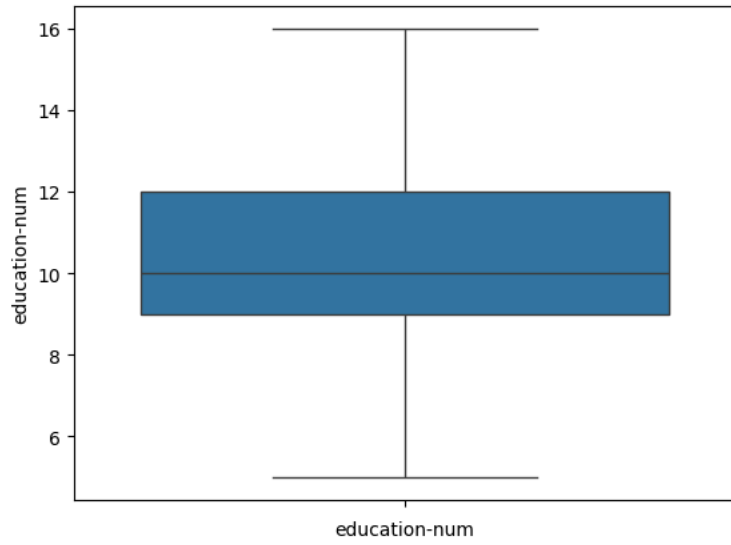





```
for x in df.columns:
    if df[x].dtypes!='object':
        q1=df[x].quantile(0.25)
        q3=df[x].quantile(0.75)
        IQR=q3-q1
        upper=q3+(1.5*IQR)
        lower=q1-(1.5*IQR)
        df=df[(df[x]>=lower) & (df[x]<=upper)]
```

```
for x in df.columns:
    if df[x].dtypes!='object':
        sns.boxplot(df[x])
        plt.xlabel(x)
        plt.show()
```



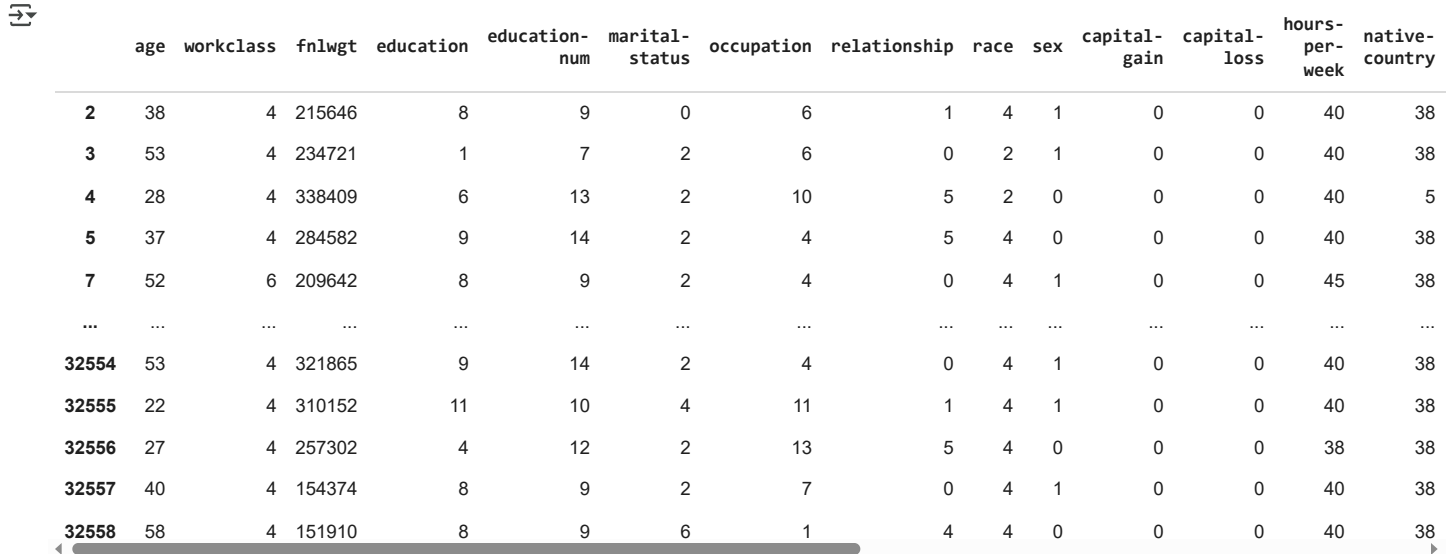


LabelEncoding

```
from sklearn.preprocessing import LabelEncoder
```

```
le=LabelEncoder()
for x in df.columns:
    if df[x].dtypes=='object':
        df[x]=le.fit_transform(df[x])
```

df



| | age | workclass | fnlwgt | education | education-num | marital-status | occupation | relationship | race | sex | capital-gain | capital-loss | hours-per-week | native-country |
|-------|-----|-----------|--------|-----------|---------------|----------------|------------|--------------|------|-----|--------------|--------------|----------------|----------------|
| 2 | 38 | 4 | 215646 | 8 | 9 | 0 | 6 | 1 | 4 | 1 | 0 | 0 | 40 | 38 |
| 3 | 53 | 4 | 234721 | 1 | 7 | 2 | 6 | 0 | 2 | 1 | 0 | 0 | 40 | 38 |
| 4 | 28 | 4 | 338409 | 6 | 13 | 2 | 10 | 5 | 2 | 0 | 0 | 0 | 40 | 5 |
| 5 | 37 | 4 | 284582 | 9 | 14 | 2 | 4 | 5 | 4 | 0 | 0 | 0 | 40 | 38 |
| 7 | 52 | 6 | 209642 | 8 | 9 | 2 | 4 | 0 | 4 | 1 | 0 | 0 | 45 | 38 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 32554 | 53 | 4 | 321865 | 9 | 14 | 2 | 4 | 0 | 4 | 1 | 0 | 0 | 40 | 38 |
| 32555 | 22 | 4 | 310152 | 11 | 10 | 4 | 11 | 1 | 4 | 1 | 0 | 0 | 40 | 38 |
| 32556 | 27 | 4 | 257302 | 4 | 12 | 2 | 13 | 5 | 4 | 0 | 0 | 0 | 38 | 38 |
| 32557 | 40 | 4 | 154374 | 8 | 9 | 2 | 7 | 0 | 4 | 1 | 0 | 0 | 40 | 38 |
| 32558 | 58 | 4 | 151910 | 8 | 9 | 6 | 1 | 4 | 4 | 0 | 0 | 0 | 40 | 38 |

Model Building


```
x=df.drop('annual_income',axis=1)
y=df['annual_income']
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score,confusion_matrix,classification_report
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=23)
```

1.Logistic Regression

```
from sklearn.linear_model import LogisticRegression
```

```
model1=LogisticRegression()
model1.fit(x_train,y_train)
```

 /usr/local/lib/python3.11/dist-packages/sklearn/linear_model/_logistic.py:465: ConvergenceWarning: lbfgs failed to converge (status=1): STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

LogisticRegression ⓘ ?

```
LogisticRegression()
```

```
y_pred=model1.predict(x_test)
```

```
accuracy_score(y_test,y_pred)*100
```


80.38957620426427

```
confusion_matrix(y_test,y_pred)
```

```
array([[2953,  78],
       [ 667, 101]])
```

```
print(classification_report(y_test,y_pred))
```

```

              precision    recall  f1-score   support

     0       0.82         0.97         0.89         3031
     1       0.56         0.13         0.21          768

 accuracy          0.80         0.80         0.80         3799
 macro avg         0.69         0.55         0.55         3799
 weighted avg      0.76         0.80         0.75         3799
```

2. Decision Trees

```
from sklearn.tree import DecisionTreeClassifier
```

```
#Brute Force to find the max_depth
```

```
max_depth=[1,2,3,4,5,6,7,8,9,10]
```

```
for i in max_depth:
```

```
    model2=DecisionTreeClassifier(max_depth=i)
```

```
    model2.fit(x_train,y_train)
```

```
    y_pred=model2.predict(x_test)
```

```
    acc=accuracy_score(y_test,y_pred)
```

```
    print(f"for the max depth {i} the accuracy score is: {acc}")
```

```

for the max depth 1 the accuracy score is: 0.7978415372466439
for the max depth 2 the accuracy score is: 0.8341668860226376
for the max depth 3 the accuracy score is: 0.8381152934982891
for the max depth 4 the accuracy score is: 0.8373256120031587
for the max depth 5 the accuracy score is: 0.8349565675177678
for the max depth 6 the accuracy score is: 0.8383785206633324
for the max depth 7 the accuracy score is: 0.8336404316925506
for the max depth 8 the accuracy score is: 0.8304817057120295
for the max depth 9 the accuracy score is: 0.8304817057120295
for the max depth 10 the accuracy score is: 0.8267965254014215
```

```
model2=DecisionTreeClassifier(max_depth=4)
```

```
model2.fit(x_train,y_train)
```

```

DecisionTreeClassifier
DecisionTreeClassifier(max_depth=4)
```

```
dt_y_pred=model2.predict(x_test)
```

```
accuracy_score(y_test,dt_y_pred)*100
```

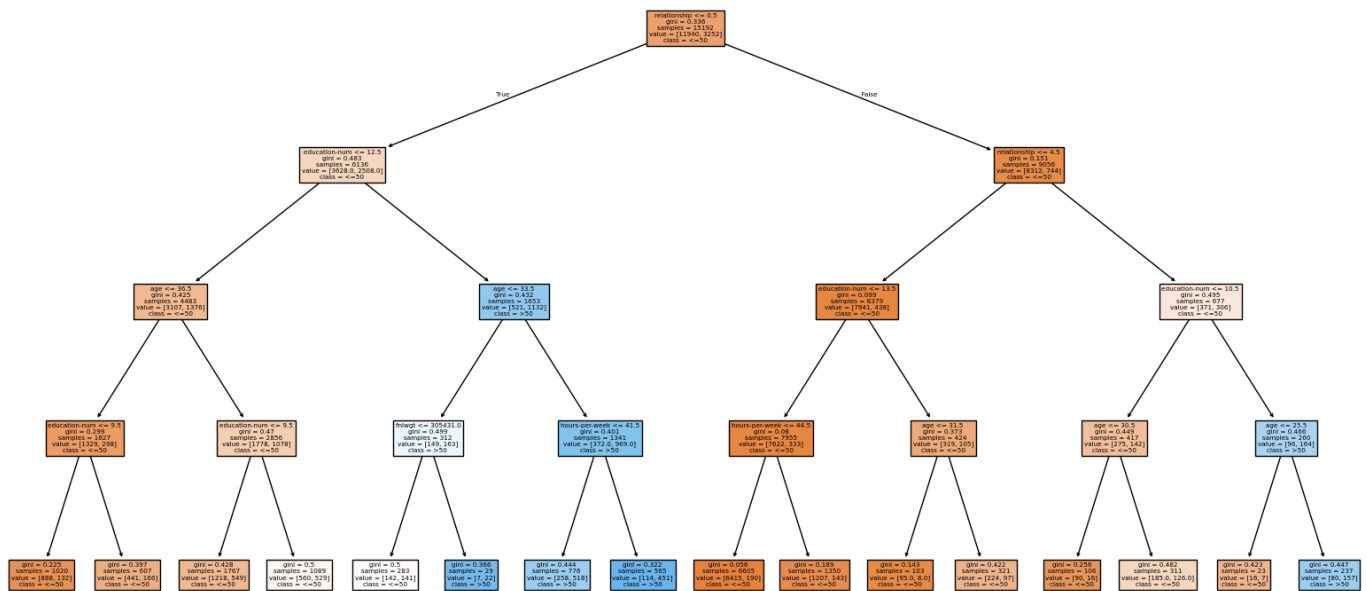
```
83.73256120031587
```

```
from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(20,10))
```

```
plot_tree(model2,filled=True,feature_names=x.columns,class_names=['<=50','>50'])
```

```
plt.show()
```



3.Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```


```
from sklearn.model_selection import RandomizedSearchCV
```

```
base_model=RandomForestClassifier(random_state=23)
```

```
params_grid={
    'n_estimators':[100,200,300],
    'max_depth':[1,5,10],
    'min_samples_split':[2,5,7],
    'min_samples_leaf':[1,2,4],
    'criterion':['gini','entropy']
}
```

```
random_search=RandomizedSearchCV(estimator=base_model,param_distributions=params_grid)
```

```
random_search.fit(x_train,y_train)
```



```
print(random_search.best_params_)
```

```
➦ {'n_estimators': 300, 'min_samples_split': 7, 'min_samples_leaf': 2, 'max_depth': 5, 'criterion': 'entropy'}
```

```
model3=RandomForestClassifier(n_estimators=300,min_samples_split=7,min_samples_leaf=2,max_depth=5,criterion='entropy',random_state=23)
```

```
model3.fit(x_train,y_train)
```

```
➦ ▼ RandomForestClassifier ⓘ ?  
RandomForestClassifier(criterion='entropy', max_depth=5, min_samples_leaf=2,  
                        min_samples_split=7, n_estimators=300, random_state=23)
```

```
rf_y_pred=model3.predict(x_test)
```

```
accuracy_score(y_test,rf_y_pred)*100
```

```
➦ 83.75888391682021
```